

EECS 570

Lecture 6

Synchronization II

Winter 2019

Prof. Thomas Wenisch

<http://www.eecs.umich.edu/courses/eecs570/>



Slides developed in part by Profs. Adve, Falsafi, Hill, Lebeck, Martin, Narayanasamy, Nowatzky, Reinhardt, Roth, Smith, Singh, and Wenisch. Some slides derived from Herlihy & Shavit "The Art of Multiprocessor Programming" used under <http://creativecommons.org/licenses/by-sa/3.0/>

Announcements

Project Proposals due - 1/30

Programming Assignment 1 due Friday 2/8 11:59pm

- Upload zip in Canvas

Readings

For Today:

- Michael Scott. *Shared-Memory Synchronization*. Morgan & Claypool Synthesis Lectures on Computer Architecture (Ch. 1, 4.0-4.3.3, 5.0-5.2.5).
- Alain Kagi, Doug Burger, and Jim Goodman. Efficient Synchronization: Let Them Eat QOLB, Proc. 24th International Symposium on Computer Architecture (ISCA 24), June, 1997.

For Wednesday:

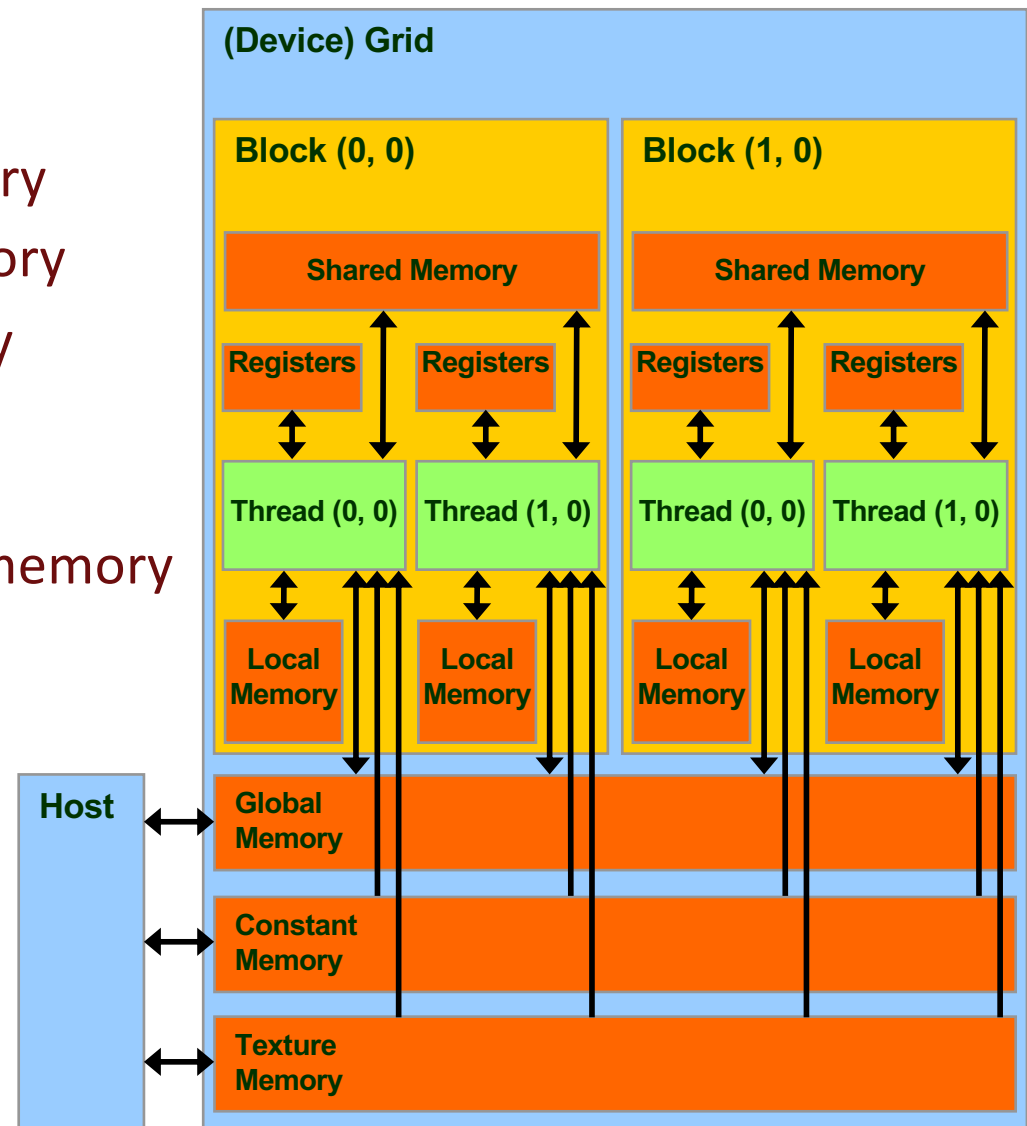
- Michael Scott. *Shared-Memory Synchronization*. Morgan & Claypool Synthesis Lectures on Computer Architecture (Ch. 8-8.3).
- M. Herlihy, Wait-Free Synchronization, ACM Trans. Program. Lang. Syst. 13(1): 124-149 (1991).

Execution Model

- Each thread block is executed by a single multiprocessor
 - Synchronized using shared memory
- Many thread blocks are assigned to a single multiprocessor
 - Executed concurrently in a time-sharing fashion
 - Keep GPU as busy as possible
- Running many threads in parallel can hide DRAM memory latency
 - Global memory access : 2~300 cycles

CUDA Device Memory Space Overview

- Each thread can:
 - R/W per-thread **registers**
 - R/W per-thread **local memory**
 - R/W per-block **shared memory**
 - R/W per-grid **global memory**
 - Read only per-grid **constant memory**
 - Read only per-grid **texture memory**
- The host can R/W **global**, **constant**, and **texture** memories



Example: Vector Addition Kernel

```
// Pair-wise addition of vector elements
// One thread per addition

__global__ void
vectorAdd(float* iA, float* iB, float* oC)
{
    int idx = threadIdx.x
        + blockDim.x * blockIdx.x;
    oC[idx] = iA[idx] + iB[idx];
}
```

Courtesy NVIDIA

Example: Vector Addition Host Code

```
float* h_A = (float*) malloc(N * sizeof(float));
float* h_B = (float*) malloc(N * sizeof(float));
// ... initialize h_A and h_B

// allocate device memory
float* d_A, d_B, d_C;
cudaMalloc( (void**) &d_A, N * sizeof(float) );
cudaMalloc( (void**) &d_B, N * sizeof(float) );
cudaMalloc( (void**) &d_C, N * sizeof(float) );

// copy host memory to device
cudaMemcpy( d_A, h_A, N * sizeof(float),
             cudaMemcpyHostToDevice );
cudaMemcpy( d_B, h_B, N * sizeof(float),
             cudaMemcpyHostToDevice );

// execute the kernel on N/256 blocks of 256 threads each
vectorAdd<<< N/256, 256>>>( d_A, d_B, d_C );
```

CUDA-Strengths

- Easy to program (small learning curve)
- Success with several complex applications
 - ▣ At least 7X faster than CPU stand-alone implementations
- Allows us to read and write data at any location in the device memory
- More fast memory close to the processors (registers + shared memory)

CUDA-Limitations

- Some hardwired graphic components are hidden
- Better tools are needed
 - ❑ Profiling
 - ❑ Memory blocking and layout
 - ❑ Binary Translation
- Difficult to find optimal values for CUDA execution parameters
 - Number of thread per block
 - Dimension and orientation of blocks and grid
 - Use of on-chip memory resources including registers and shared memory

Synchronization

Synchronization objectives

- Low overhead
 - Synchronization can limit scalability (E.g., single-lock OS kernels)
- Correctness (and ease of programmability)
 - Synchronization failures are extremely difficult to debug
- Coordination of HW and SW
 - SW semantics must be tightly specified to prove correctness
 - HW can often improve efficiency

Synchronization Forms

- Mutual exclusion (critical sections)
 - Lock & Unlock
- Event Notification
 - Point-to-point (producer-consumer, flags)
 - I/O, interrupts, exceptions
- Barrier Synchronization
- Higher-level constructs
 - Queues, software pipelines, (virtual) time, counters
- Next lecture: optimistic concurrency control
 - Transactional Memory

Anatomy of a Synchronization Op

- Acquire Method
 - Way to obtain the lock or proceed past the barrier
- Waiting Algorithm
 - Spin (aka busy wait)
 - Waiting process repeatedly tests a location until it changes
 - Releasing process sets the location
 - Lower overhead, but wastes CPU resources
 - Can cause interconnect traffic
 - Block (aka suspend)
 - Waiting process is descheduled
 - High overhead, but frees CPU to do other things
 - Hybrids (e.g., spin, then block)
- Release Method
 - Way to allow other processes to proceed

HW/SW Implementation Trade-offs

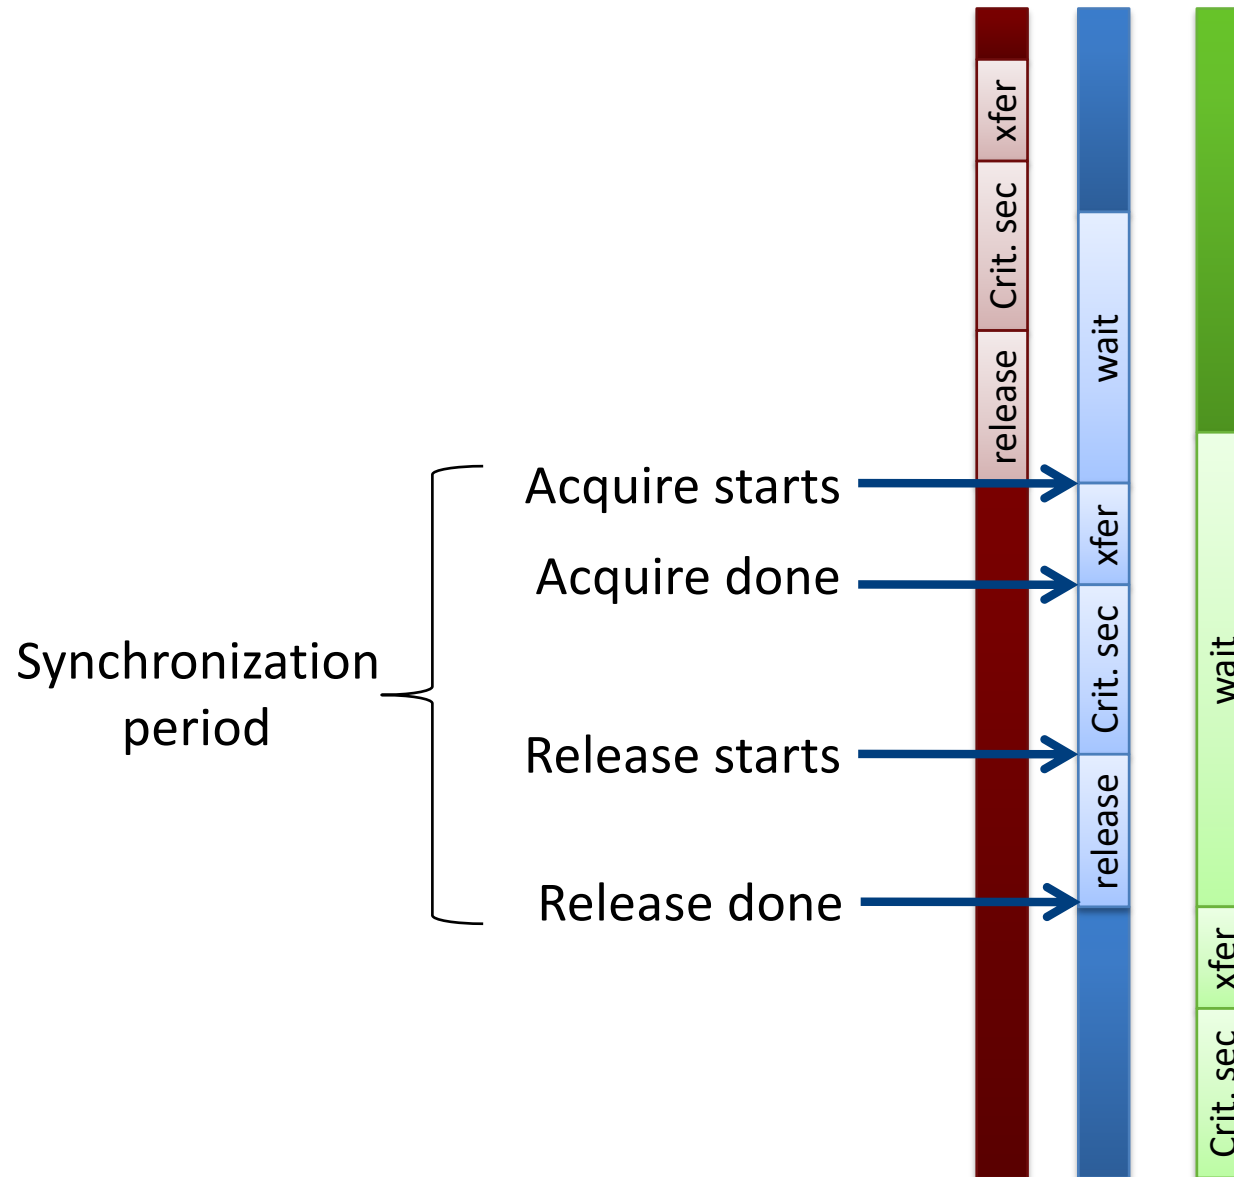
- User wants high-level (ease of programming)
 - ❑ LOCK(lock_variable); UNLOCK(lock_variable)
 - ❑ BARRIER(barrier_variable, numprocs)
- SW advantages: flexibility, portability
- HW advantages: speed
- Design objectives:
 - ❑ Low latency
 - ❑ Low traffic
 - ❑ Low storage
 - ❑ Scalability (“wait-free”-ness)
 - ❑ Fairness

Challenges

- Same sync may have different behavior at different times
 - ❑ Lock accessed with low or high contention
 - ❑ Different performance needs: low latency vs. high throughput
 - ❑ Different algorithms best for each, need different primitives
- Multiprogramming can change sync behavior
 - ❑ Process scheduling or other resource interactions
 - ❑ May need algorithms that are worse in dedicated case
- Rich area of SW/HW interactions
 - ❑ Which primitives are available?
 - ❑ What communication patterns cost more/less?

Locks

Lock-based Mutual Exclusion



No contention:

- Want low latency

Contention:

- Want low period
- Low traffic
- Fairness

How Not to Implement Locks

- LOCK

```
while (lock_variable == 1);
```

```
lock_variable = 1;
```



Context switch!

- UNLOCK

```
lock_variable = 0;
```

Solution: Atomic Read-Modify-Write

- Test&Set(r,x)
`{ r=m[x] ; m[x]=1 ; }`
 - r is register
 - m[x] is memory location x
- Fetch&Op(r1,r2,x,op)
`{ r1=m[x] ; m[x]=op(r1,r2) ; }`
- Swap(r,x)
`{ temp=m[x] ; m[x]=r ; r=temp ; }`
- Compare&Swap(r1,r2,x)
`{ temp=r2 ; r2=m[x] ; if r1==r2 then m[x]=temp ; }`

Implementing RMWs

- Bus-based systems:
 - Hold bus and issue load/store operations without any intervening accesses by other processors
- Scalable systems
 - Acquire exclusive ownership via cache coherence
 - Perform load/store operations without allowing external coherence requests

Load-Locked Store-Conditional

- Load-locked
 - ❑ Issues a normal load...
 - ❑ ...and sets a **flag** and **address** field
- Store-conditional
 - ❑ Checks that **flag** is set and **address** matches...
 - ❑ ...only then performs store
- **Flag** is cleared by
 - ❑ Invalidation
 - ❑ Cache eviction
 - ❑ Context switch

```
lock: while (1) {  
    load-locked r1, lock_variable  
    if (r1 == 0) {  
        mov r2 = 1  
        if (SC r2, lock) break;  
    }  
}  
  
unlock: st lock_variable, #0
```

Test-and-Set Spin Lock (T&S)

- Lock is “acquire”, Unlock is “release”
- `acquire(lock_ptr) :`
 - `while (true) :`
 - // Perform “test-and-set”*
 - `old = compare_and_swap(lock_ptr, UNLOCKED, LOCKED)`
 - `if (old == UNLOCKED) :`
 - `break // lock acquired!`
 - // keep spinning, back to top of while loop*
- `release(lock_ptr) :`
 - `store[lock_ptr] <- UNLOCKED`
- Performance problem
 - ❑ CAS is both a read and write; spinning causes lots of invalidations

Test-and-Test-and-Set Spin Lock (TTS)

- `acquire(lock_ptr) :`
 - `while (true) :`
 - // Perform “test”*
 - `load [lock_ptr] -> original_value`
 - `if (original_value == UNLOCKED) :`
 - // Perform “test-and-set”*
 - `old = compare_and_swap(lock_ptr, UNLOCKED, LOCKED)`
 - `if (old == UNLOCKED) :`
 - `break // lock acquired!`
 - // keep spinning, back to top of while loop*
- `release(lock_ptr) :`
 - `store[lock_ptr] <- UNLOCKED`
- **Now “spinning” is read-only, on local cached copy**

TTS Lock Performance Issues

- **Performance issues remain**

- ☐ Every time the lock is released...
- ☐ All the processors load it, and likely try to CAS the block
- ☐ Causes a storm of coherence traffic, clogs things up badly

- **One solution: backoff**

- ☐ Instead of spinning constantly, check less frequently
- ☐ Exponential backoff works well in practice

- **Another problem with spinning**

- ☐ Processors can spin really fast, starve threads on the same core!
- ☐ Solution: x86 adds a “PAUSE” instruction
 - Tells processor to suspend the thread for a short time

- **(Un)fairness**

Ticket Locks

- **To ensure fairness and reduce coherence storms**
- Locks have two counters: `next_ticket`, `now_serving`
 - Deli counter
- `acquire(lock_ptr) :`
 - `my_ticket = fetch_and_increment(lock_ptr->next_ticket)`
 - `while(lock_ptr->now_serving != my_ticket); // spin`
- `release(lock_ptr) :`
 - `lock_ptr->now_serving = lock_ptr->now_serving + 1`
 - (Just a normal store, not an atomic operation, why?)
- Summary of operation
 - To “get in line” to acquire the lock, CAS on `next_ticket`
 - Spin on `now_serving`

Ticket Locks

- **Properties**

- ❑ Less of a “thundering herd” coherence storm problem
 - To acquire, only need to read new value of `now_serving`
- ❑ No CAS on critical path of lock handoff
 - Just a non-atomic store
- ❑ FIFO order (fair)
 - Good, but only if the O.S. hasn’t swapped out any threads!

- **Padding**

- ❑ Allocate `now_serving` and `next_ticket` on different cache blocks
 - `struct { int now_serving; char pad[60]; int next_ticket; } ...`
- ❑ Two locations reduces interference

- **Proportional backoff**

- ❑ Estimate of wait time: $(\text{my_ticket} - \text{now_serving}) * \text{average hold time}$

Array-Based Queue Locks

- **Why not give each waiter its own location to spin on?**
 - ❑ Avoid coherence storms altogether!
- **Idea: “slot” array of size N: “go ahead” or “must wait”**
 - Initialize first slot to “go ahead”, all others to “must wait”
 - Padded one slot per cache block,
 - ❑ Keep a “next slot” counter (similar to “next_ticket” counter)
- Acquire: “get in line”
 - ❑ $\text{my_slot} = (\text{atomic increment of “next slot” counter}) \bmod N$
 - ❑ Spin while $\text{slots}[\text{my_slot}]$ contains “must_wait”
 - ❑ Reset $\text{slots}[\text{my_slot}]$ to “must wait”
- Release: “unblock next in line”
 - ❑ Set $\text{slots}[\text{my_slot} + 1 \bmod N]$ to “go ahead”

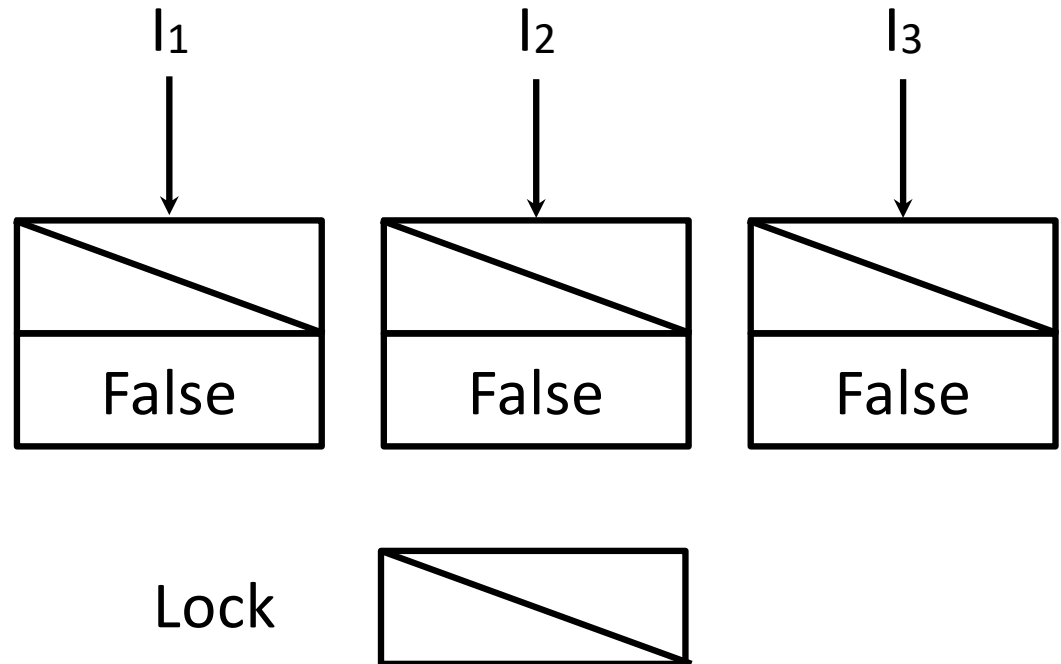
Array-Based Queue Locks

- Variants: Anderson 1990, Graunke and Thakkar 1990
- Desirable properties
 - ❑ Threads spin on dedicated location
 - Just two coherence misses per handoff
 - Traffic independent of number of waiters
 - ❑ FIFO & fair (same as ticket lock)
- Undesirable properties
 - ❑ Higher uncontended overhead than a TTS lock
 - ❑ Storage $O(N)$ for each lock
 - 128 threads at 64B padding: 8KBs per lock!
 - What if N isn't known at start?
- List-based locks address the $O(N)$ storage problem
 - ❑ Several variants of list-based locks: MCS 1991, CLH 1993/1994

List-Based Queue Lock (MCS)

- A “lock” is a pointer to a linked list node
 - ❑ next node pointer
 - ❑ boolean must_wait
 - ❑ Each thread has its own local pointer to a node “I”
- `acquire(lock) :`
 - `I->next = null;`
 - `predecessor = fetch_and_store(lock, I)`
 - `if predecessor != nil` *//some node holds lock*
 - `I->must_wait = true`
 - `predecessor->next = I` *//predecessor must wake us*
 - `repeat while I->must_wait` *//spin till lock is free*
- `release(lock) :`
 - `if (I->next == null)` *//no known successor*
 - `if compare_and_swap(lock, I, nil)` *//make sure...*
 - `return` *//CAS succeeded; lock freed*
 - `repeat while I->next = nil` *//spin to learn successor*
 - `I->next->must_wait = false` *//wake successor*

MCS Lock Example: Time 0

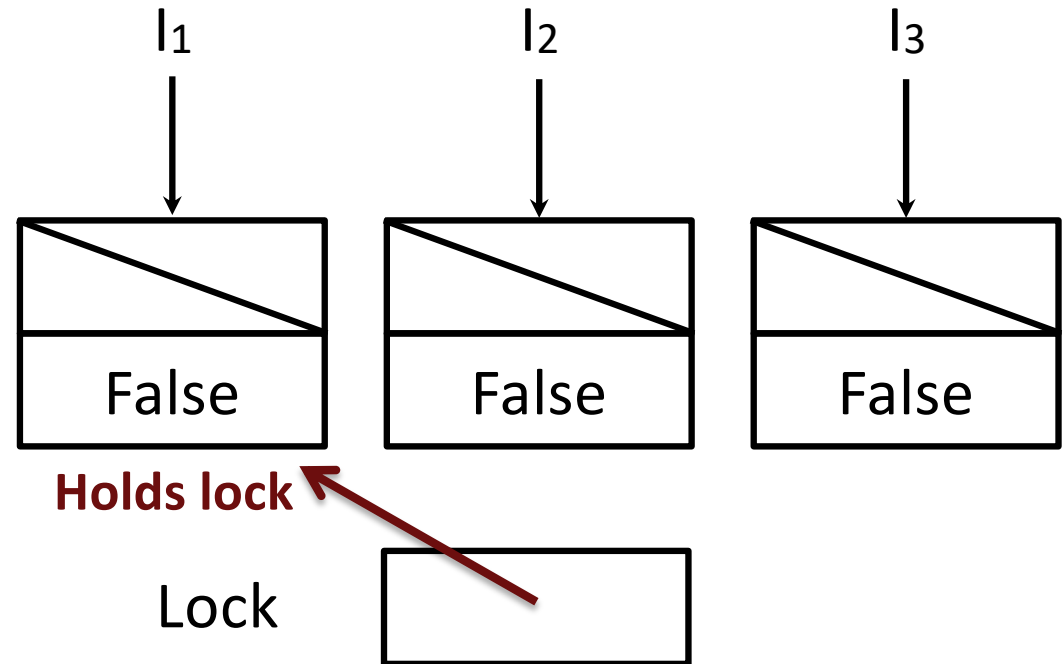


- `acquire(lock) :`
 `I->next = null;`
 `pred = FAS(lock, I)`
 `if pred != nil`
 `I->must_wait = true`
 `pred->next = I`
 `repeat while I->must_wait`

- `release(lock) :`
 `if (I->next == null)`
 `if CAS(lock, I, nil)`
 `return`
 `repeat while I->next == nil`
 `I->next->must_wait = false`

MCS Lock Example: Time 1

- t_1 : Acquire(L)

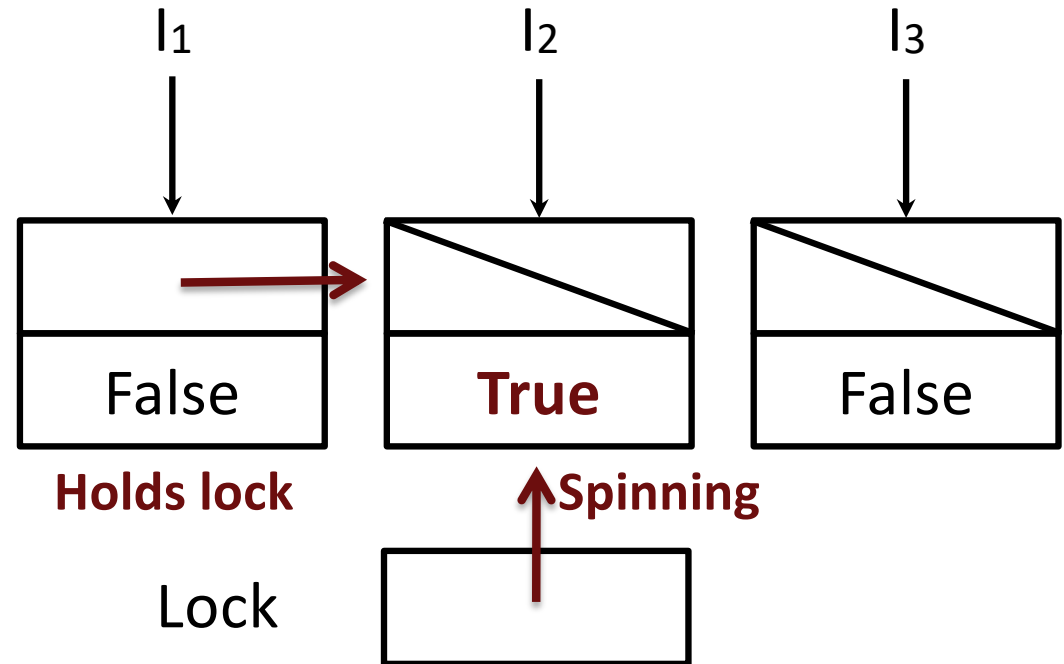


```
• acquire(lock) :  
  I->next = null;  
  pred = FAS(lock, I)  
  if pred != nil  
    I->must_wait = true  
  pred->next = I  
  repeat while I->must_wait
```

```
• release(lock) :  
  if (I->next == null)  
    if CAS(lock, I, nil)  
      return  
  repeat while I->next == nil  
  I->next->must_wait = false
```

MCS Lock Example: Time 2

- t_1 : Acquire(L)
- t_2 : Acquire(L)



```

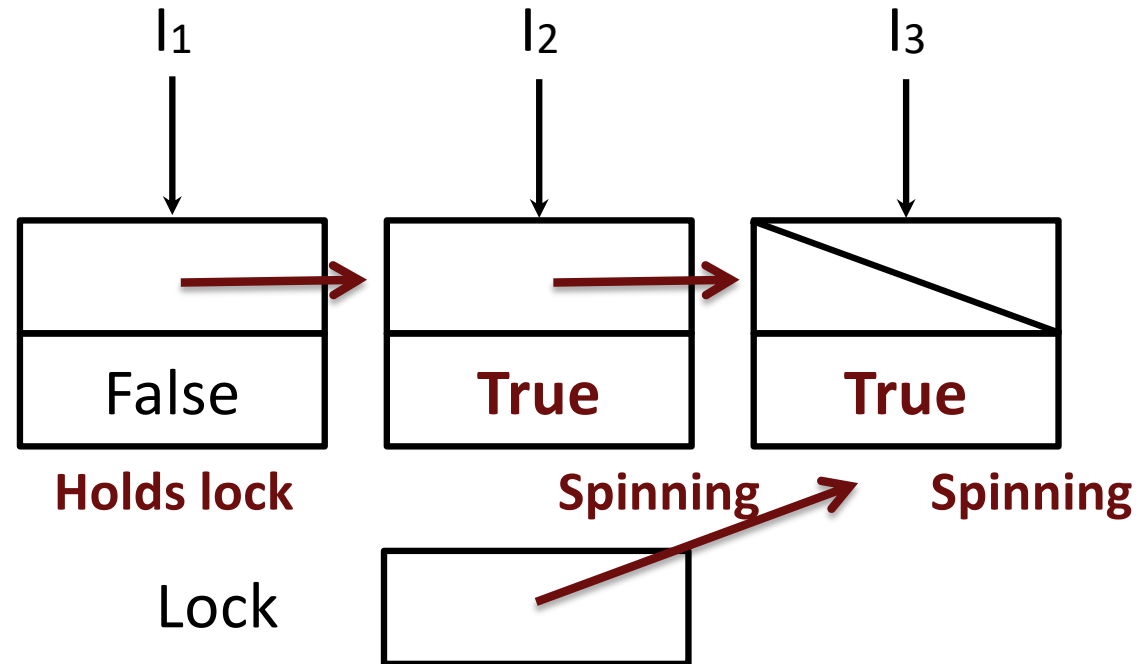
• acquire(lock) :
  I->next = null;
  pred = FAS(lock, I)
  if pred != nil
    I->must_wait = true
  pred->next = I
  repeat while I->must_wait
  
```

```

• release(lock) :
  if (I->next == null)
    if CAS(lock, I, nil)
      return
  repeat while I->next == nil
  I->next->must_wait = false
  
```


MCS Lock Example: Time 3

- t_1 : Acquire(L)
- t_2 : Acquire(L)
- t_3 : Acquire(L)



```

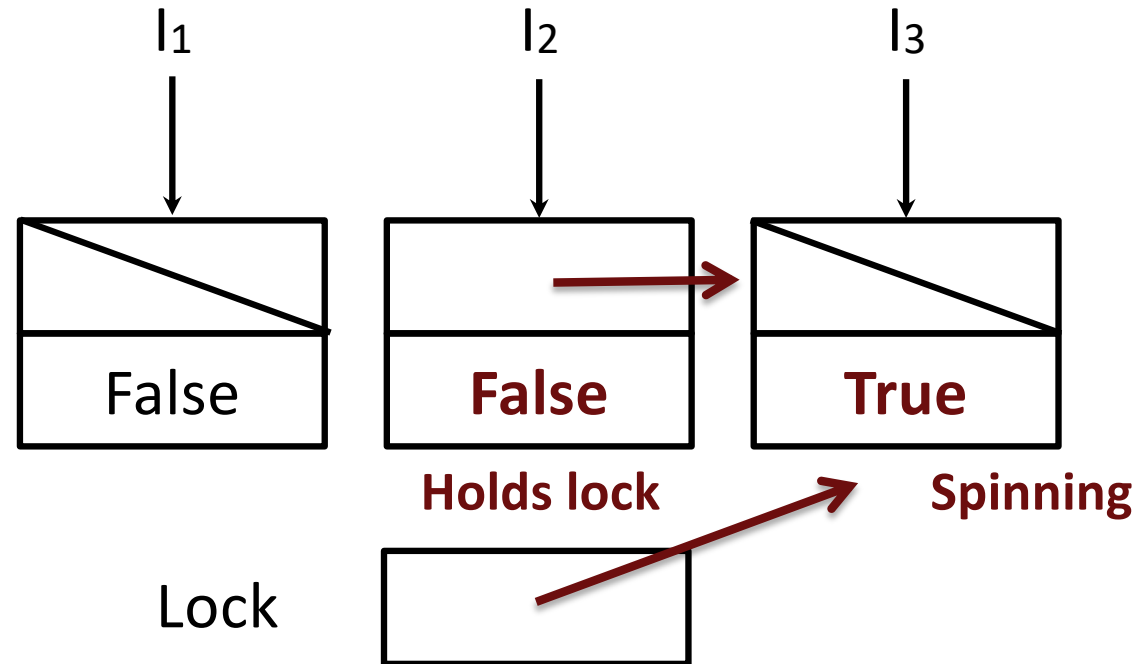
• acquire(lock) :
  I->next = null;
  pred = FAS(lock, I)
  if pred != nil
    I->must_wait = true
  pred->next = I
  repeat while I->must_wait
  
```

```

• release(lock) :
  if (I->next == null)
    if CAS(lock, I, nil)
      return
  repeat while I->next == nil
  I->next->must_wait = false
  
```

MCS Lock Example: Time 4

- t₁: Acquire(L)
- t₂: Acquire(L)
- t₃: Acquire(L)
- t₁: Release(L)



```

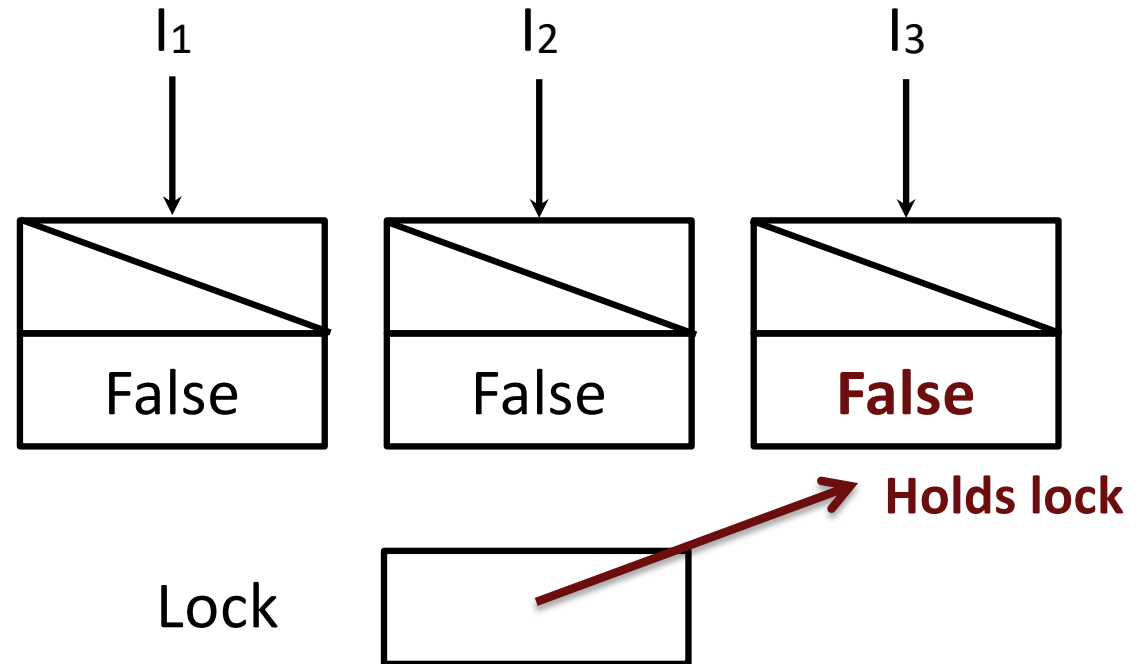
• acquire(lock) :
  I->next = null;
  pred = FAS(lock, I)
  if pred != nil
    I->must_wait = true
    pred->next = I
  repeat while I->must_wait
  
```

```

• release(lock) :
  if (I->next == null)
    if CAS(lock, I, nil)
      return
  repeat while I->next == nil
  I->next->must_wait = false
  
```

MCS Lock Example: Time 5

- t_1 : Acquire(L)
- t_2 : Acquire(L)
- t_3 : Acquire(L)
- t_1 : Release(L)
- t_2 : Release(L)



```

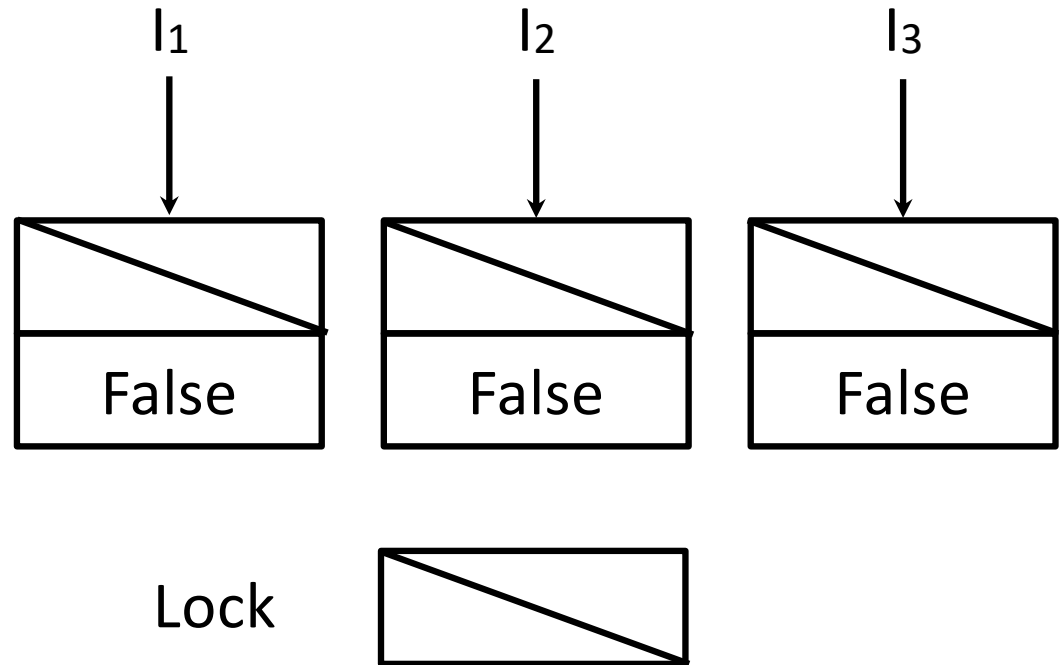
• acquire(lock) :
  I->next = null;
  pred = FAS(lock, I)
  if pred != nil
    I->must_wait = true
  pred->next = I
  repeat while I->must_wait
  
```

```

• release(lock) :
  if (I->next == null)
    if CAS(lock, I, nil)
      return
  repeat while I->next == nil
  I->next->must_wait = false
  
```

MCS Lock Example: Time 6

- t_1 : Acquire(L)
- t_2 : Acquire(L)
- t_3 : Acquire(L)
- t_1 : Release(L)
- t_2 : Release(L)
- t_3 : Release(L)



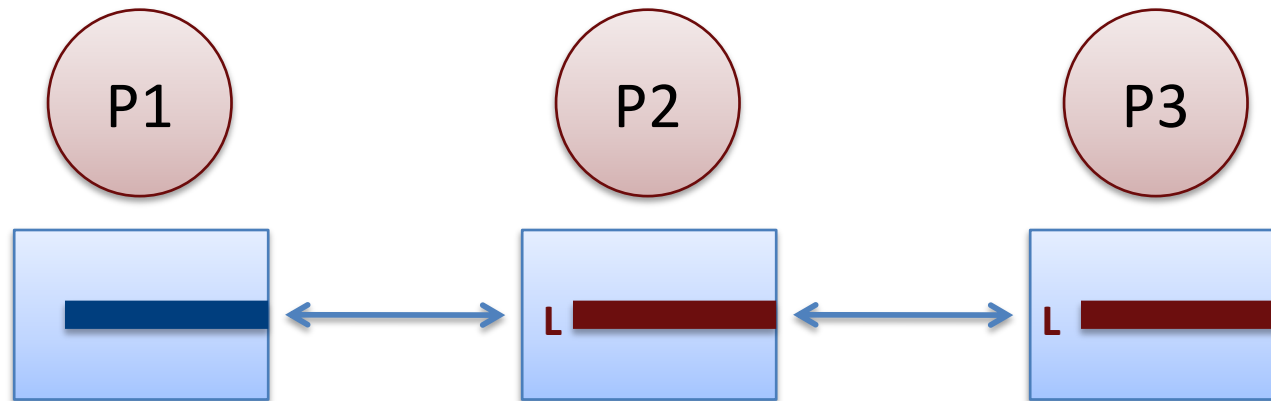
```
• acquire(lock) :  
  I->next = null;  
  pred = FAS(lock, I)  
  if pred != nil  
    I->must_wait = true  
  pred->next = I  
  repeat while I->must_wait
```

```
• release(lock) :  
  if (I->next == null)  
    if CAS(lock, I, nil)  
      return  
  repeat while I->next == nil  
  I->next->must_wait = false
```

Queue-based locks in HW: QOLB

- **Queue On Lock Bit**

- ❑ HW maintains doubly-linked list between requesters
 - This is a key idea of “Scalable Coherence Interface”, see Unit 3
- ❑ Augment cache with “locked” bit
 - Waiting caches spin on local “locked” cache line
- ❑ Upon release, lock holder sends line to 1st requester
 - Only requires one message on interconnect



Fundamental Mechanisms to Reduce Overheads

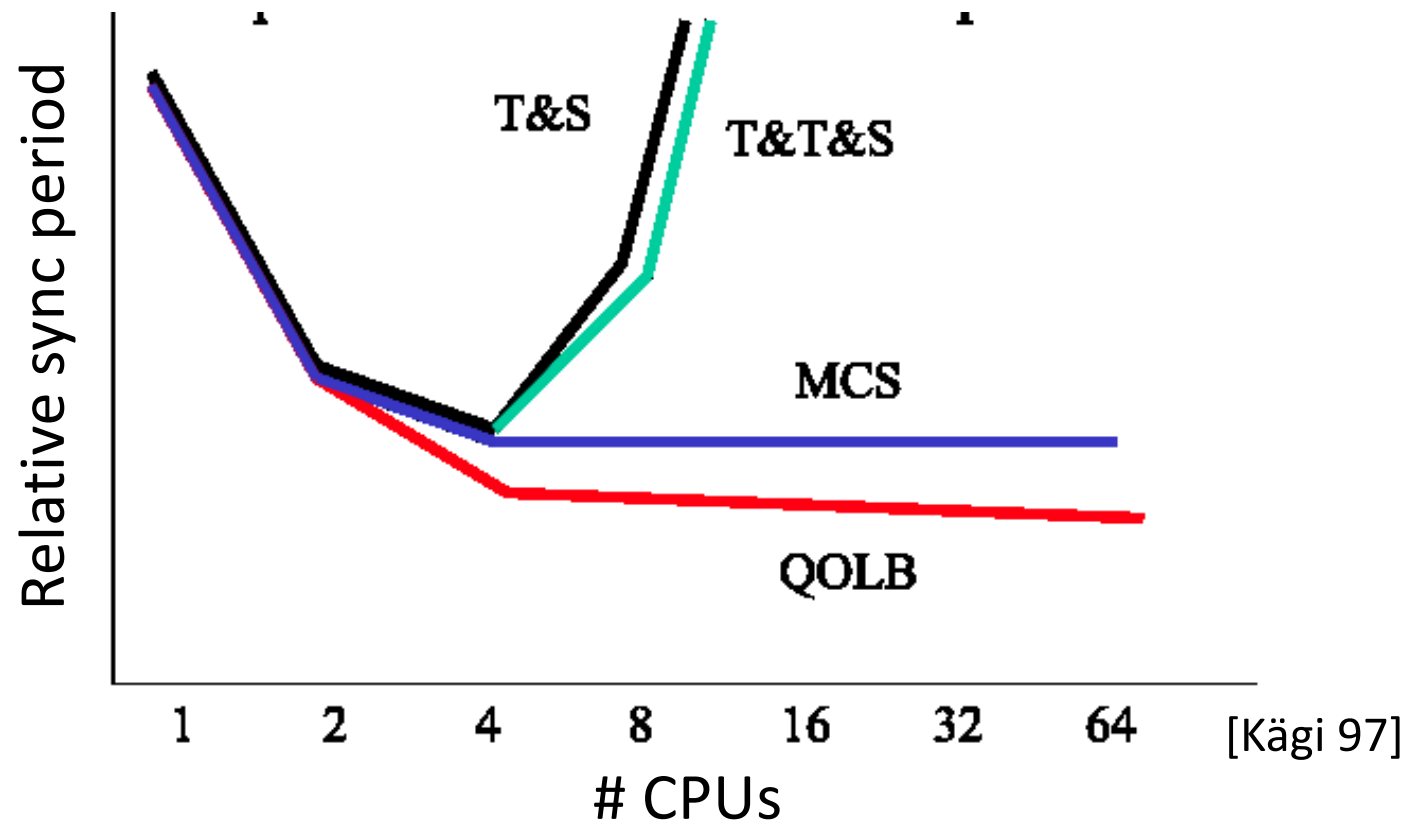
[Kägi, Burger, Goodman ASPLOS 97]

- **Basic mechanisms**

- ☐ Local Spinning
- ☐ Queue-based locking
- ☐ Collocation
- ☐ Synchronous Prefetch

	Local Spin	Queue	Collocation	Prefetch
T&S	No	No	Optional	No
T&T&S	Yes	No	Optional	No
MCS	Yes	Yes	Partial	No
QOLB	yes	Yes	Optional	Yes

Microbenchmark Analysis



Performance of Locks

- **Contention vs. No Contention**

- ☐ Test-and-Set best when no contention
- ☐ Queue-based is best with medium contention
- ☐ Idea: switch implementation based on lock behavior
 - Reactive Synchronization – Lim & Agarwal 1994
 - SmartLocks – Eastep et al 2009

- **High-contention indicates poorly written program**

- ☐ Need better algorithm or data structures

Point-to-Point Event Synchronization

- Can use normal variables as flags

```
a = f(x);
```

```
flag = 1;
```

```
while (flag == 0);
```

```
b = g(a);
```

- If we know initial conditions

```
a = f(x);
```

```
while (a == 0);
```

```
b = g(a);
```

- **Assumes Sequential Consistency!**

- Full/Empty Bits

- ☐ Set on write

- ☐ Cleared on read

- ☐ Can't write if set, can't read if clear

Barriers

Barriers

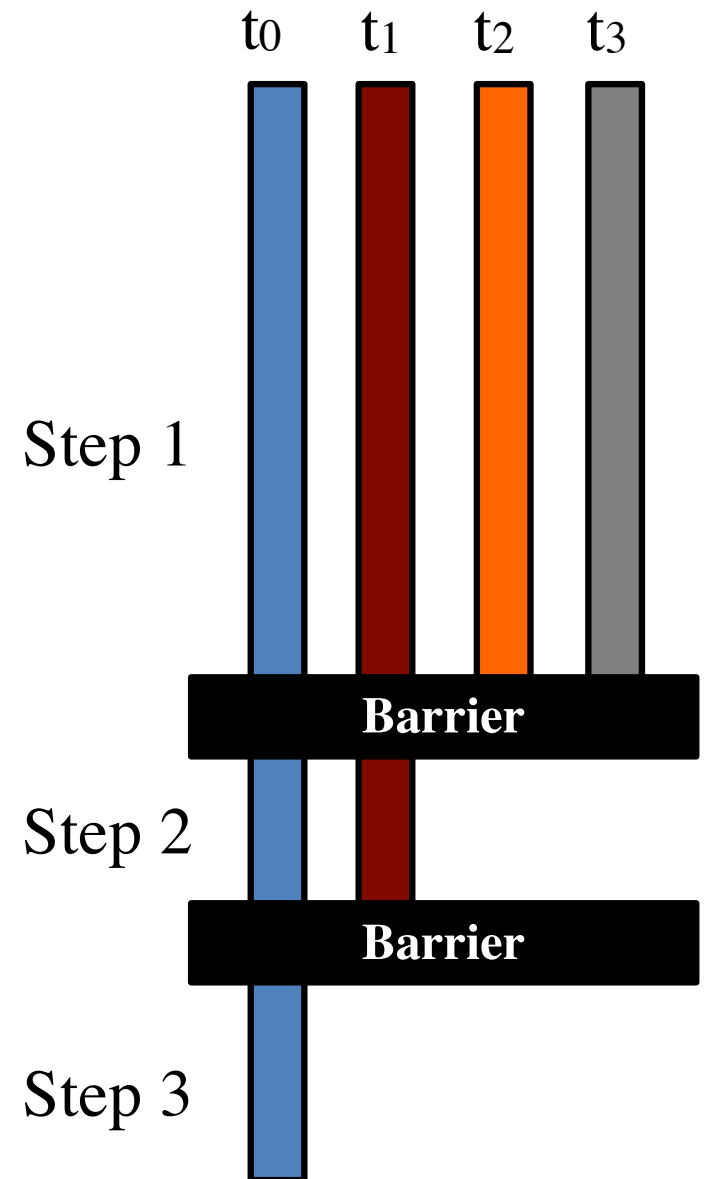
- Physics simulation computation
 - ❑ Divide up each timestep computation into N independent pieces
 - ❑ Each timestep: compute independently, synchronize

- Example: each thread executes:

```
segment_size = total_particles / number_of_threads
my_start_particle = thread_id * segment_size
my_end_particle = my_start_particle + segment_size - 1
for (timestep = 0; timestep += delta; timestep < stop_time):
    calculate_forces(t, my_start_particle, my_end_particle)
    barrier()
    update_locations(t, my_start_particle, my_end_particle)
    barrier()
```

- Barrier? All threads wait until all threads have reached it

Example: Barrier-Based Merge Sort



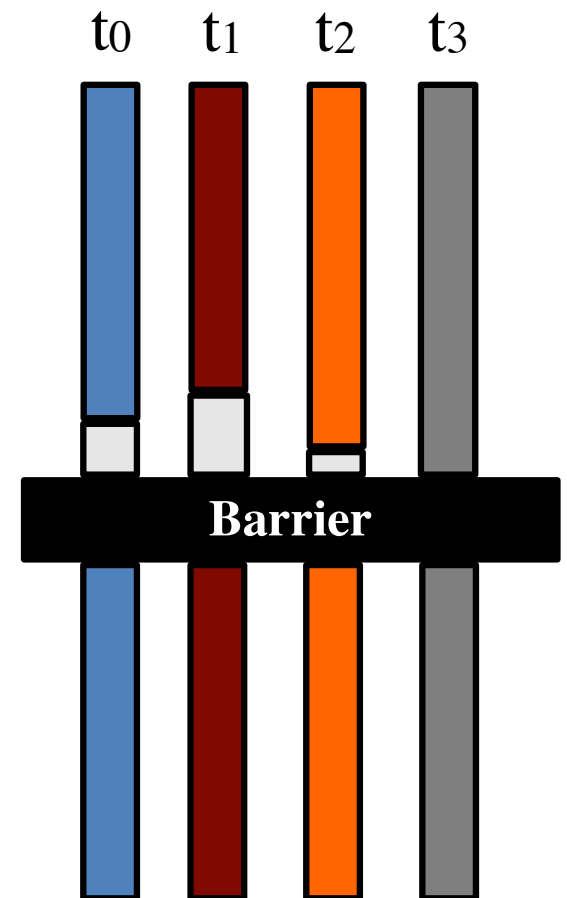
Global Synchronization Barrier

- At a barrier
 - ❑ All threads wait until all other threads have reached it
- Strawman implementation (**wrong!**)

```
global (shared) count : integer := P
```

```
procedure central_barrier
  if fetch_and_decrement(&count) == 1
    count := P
  else
    repeat until count == P
```

- What is wrong with the above code?



Sense-Reversing Barriers

```
public class Barrier {
    AtomicInteger count;
    int size;
    boolean sense = false;
    ThreadLocal<boolean> threadSense = new ThreadLocal<boolean>();

    public void await () {
        boolean mySense = threadSense.get();
        if (count.getAndDecrement()==1) {
            count.set(size); sense = !mySense;
        } else {
            while (sense != mySense) {}
        }
        threadSense.set(!mySense);
    }
}
```

Sense-Reversing Barriers

```
public class Barrier {  
    AtomicInteger count;  
    int size;  
    boolean sense = false;  
    ThreadLocal<boolean> threadSense = new ThreadLocal<boolean>();  
  
    public void await {  
        boolean mySense = threadSense.get();  
        if (count.getAndDecrement()==1) {  
            count.set(size); sense = !mySense  
        } else {  
            while (sense != mySense) {}  
        }  
        threadSense.set(!mySense)}}}
```

Completed odd or
even-numbered
phase?

Sense-Reversing Barriers

```
public class Barrier {  
    AtomicInteger count;  
    int size;  
    boolean sense = false;
```

**Store sense for
next phase**

```
    threadSense = new ThreadLocal<boolean>...
```

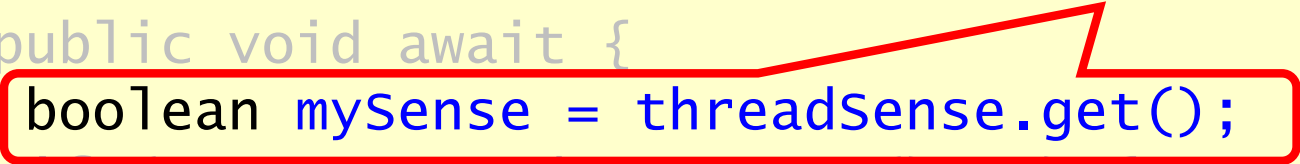
```
public void await {  
    boolean mySense = threadSense.get();  
    if (count.getAndDecrement()==1) {  
        count.set(size); sense = !mySense  
    } else {  
        while (sense != mySense) {}  
    }  
    threadSense.set(!mySense)}}}
```


Sense-Reversing Barriers

```
public class Barrier {
    AtomicInteger count;
    int size;
    boolean sense = false;
    ThreadLocal<boolean> threadSense = new ThreadLocal<boolean>();

    public void await {
        boolean mySense = threadSense.get();
        if (count.getAndDecrement() == 1) {
            count.set(size); sense = !mySense;
        } else {
            while (sense != mySense) {}
        }
        threadSense.set(!mySense);
    }
}
```

**Get new sense
determined by last
phase**



Sense-Reversing Barriers

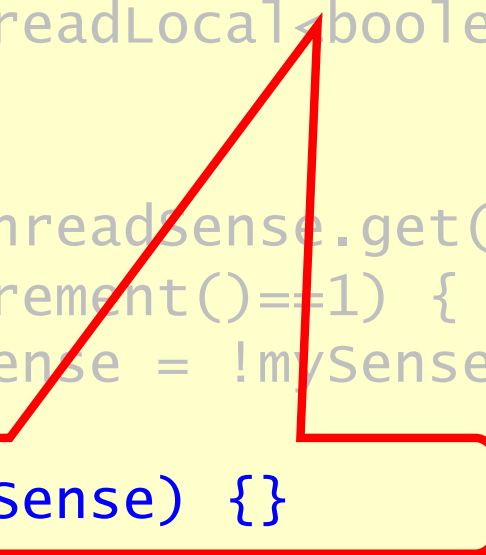
```
public class Barrier {  
    AtomicInteger count;  
    int size;  
    boolean sense = false;  
    ThreadLocal<boolean> threadSense = new ThreadLocal<boolean>();  
  
    public void await {  
        boolean mySense = threadSense.get();  
        if (count.getAndDecrement()==1) {  
            count.set(size); sense = !mySense  
        } else {  
            while (sense != mySense) {}  
        }  
        threadSense.set(!mySense)}}}
```

If I'm last, reverse sense for next time

Sense-Reversing Barriers

```
public class Barrier {  
    AtomicInteger count;  
    int size;  
    boolean sense = false;  
    ThreadLocal<boolean> threadSense = new ThreadLocal<boolean>();  
  
    public void await {  
        boolean mySense = threadSense.get();  
        if (count.getAndDecrement() == 1) {  
            count.set(size); sense = !mySense;  
        } else {  
            while (sense != mySense) {}  
        }  
        threadSense.set(!mySense);  
    }  
}
```

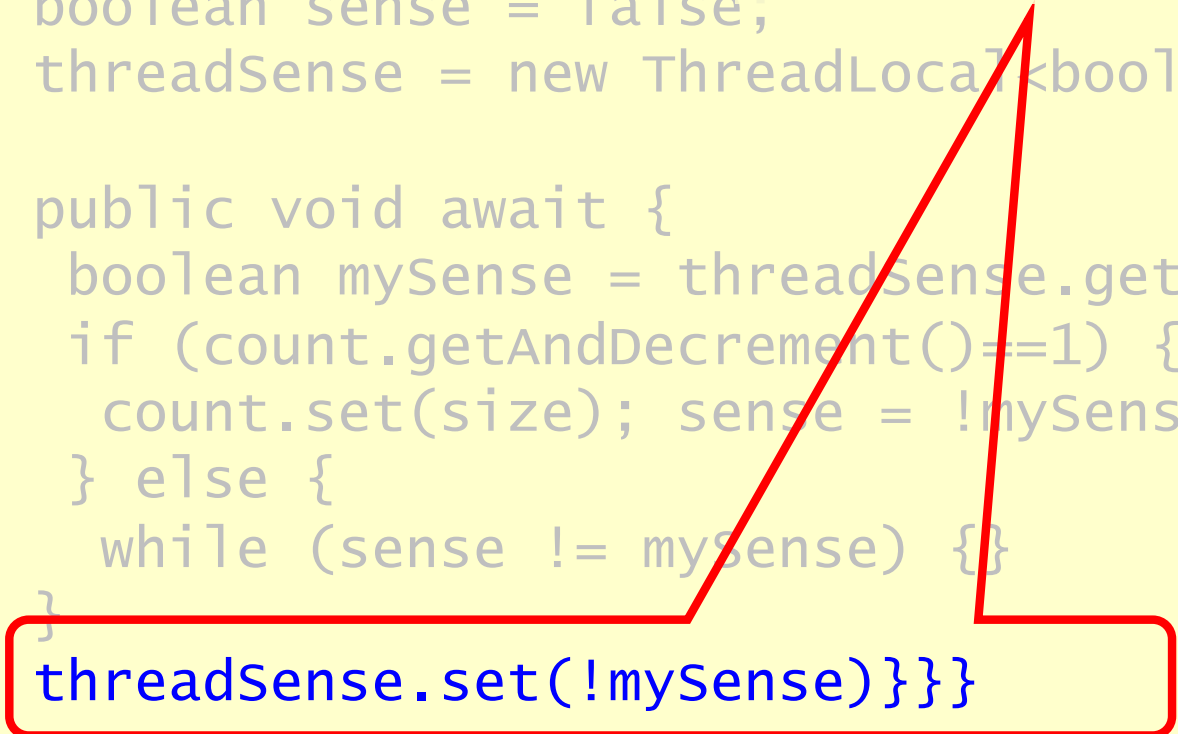
Otherwise, wait for sense to flip



Sense-Reversing Barriers

```
public class Barrier {  
    AtomicInteger count;  
    int size;  
    boolean sense = false;  
    ThreadLocal<boolean> threadSense = new ThreadLocal<boolean>();  
  
    public void await {  
        boolean mySense = threadSense.get();  
        if (count.getAndDecrement() == 1) {  
            count.set(size); sense = !mySense;  
        } else {  
            while (sense != mySense) {}  
        }  
        threadSense.set(!mySense);  
    }  
}
```

Prepare sense for next phase



Other Barrier Implementations

- Problem with centralized barrier
 - ❑ All processors must increment each counter
 - ❑ Each read/modify/write is a serialized coherence action
 - Each one is a cache miss
 - ❑ $O(n)$ if threads arrive simultaneously, slow for lots of processors
- Combining Tree Barrier
 - ❑ Build a $\log_k(n)$ height tree of counters (one per cache block)
 - ❑ Each thread coordinates with k other threads (by thread id)
 - ❑ Last of the k processors, coordinates with next higher node in tree
 - ❑ As many coordination address are used, misses are not serialized
 - ❑ $O(\log n)$ in best case
- Static and more dynamic variants
 - ❑ Tree-based arrival, tree-based or centralized release

Combining Tree Barrier

```
public class Node{
    AtomicInteger count; int size;
    Node parent; volatile boolean sense;

    public void await() {...
        if (count.getAndDecrement()==1) {
            if (parent != null) {
                parent.await()}
            count.set(size);
            sense = mySense
        } else {
            while (sense != mySense) {}
        }
    }...}}}
```

Combining Tree Barrier

```
public class Node{
    AtomicInteger count; int size;
    Node parent; volatile boolean sense;

    public void await() {...
        if (count.getAndDecrement()==1) {
            if (parent != null) {
                parent.await();
            }
            count.set(size);
            sense = mySense;
        } else {
            while (sense != mySense) {}
        }
    }...}}}
```

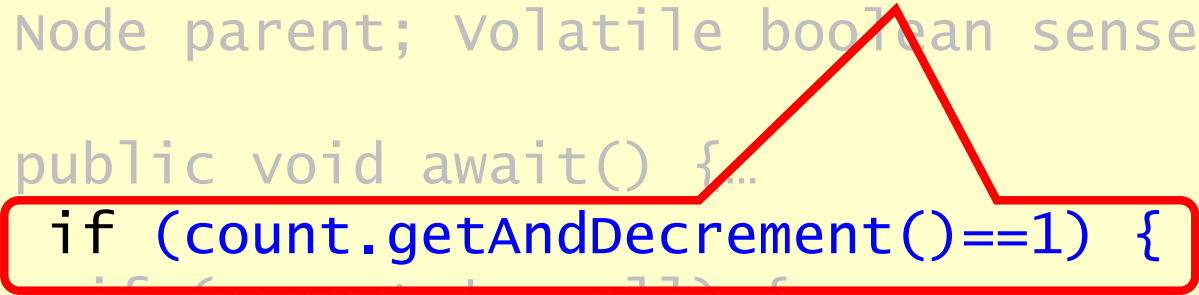
Parent barrier in tree

Combining Tree Barrier

```
public class Node{
    AtomicInteger count; int size;
    Node parent; volatile boolean sense;

    public void await() {...
        if (count.getAndDecrement()==1) {
            if (parent != null) {
                parent.await();
            }
            count.set(size);
            sense = mySense;
        } else {
            while (sense != mySense) {}
        }
    }...}}}
```

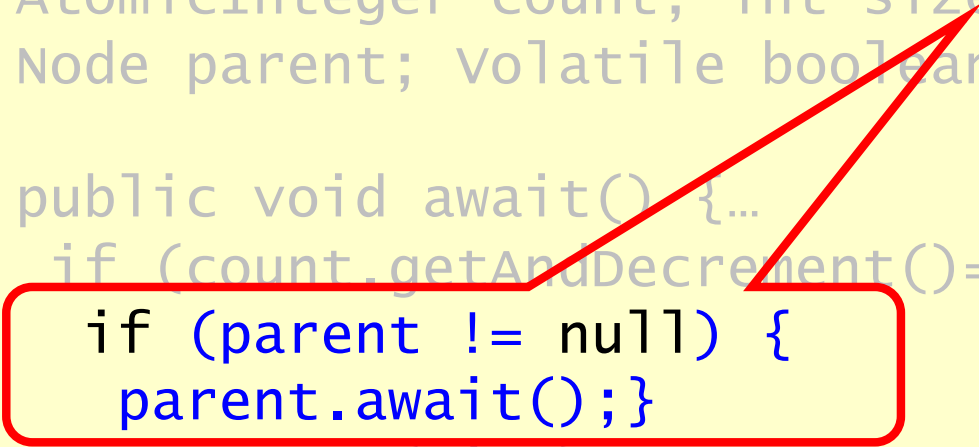
Am I last?



Combining Tree Barrier

```
public class Node{ Proceed to parent barrier
    AtomicInteger count; int size;
    Node parent; volatile boolean sense;

    public void await() {...
        if (count.getAndDecrement()==1) {
            if (parent != null) {
                parent.await();}
            count.set(size);
            sense = mySense
        } else {
            while (sense != mySense) {}
        }...}}}
```

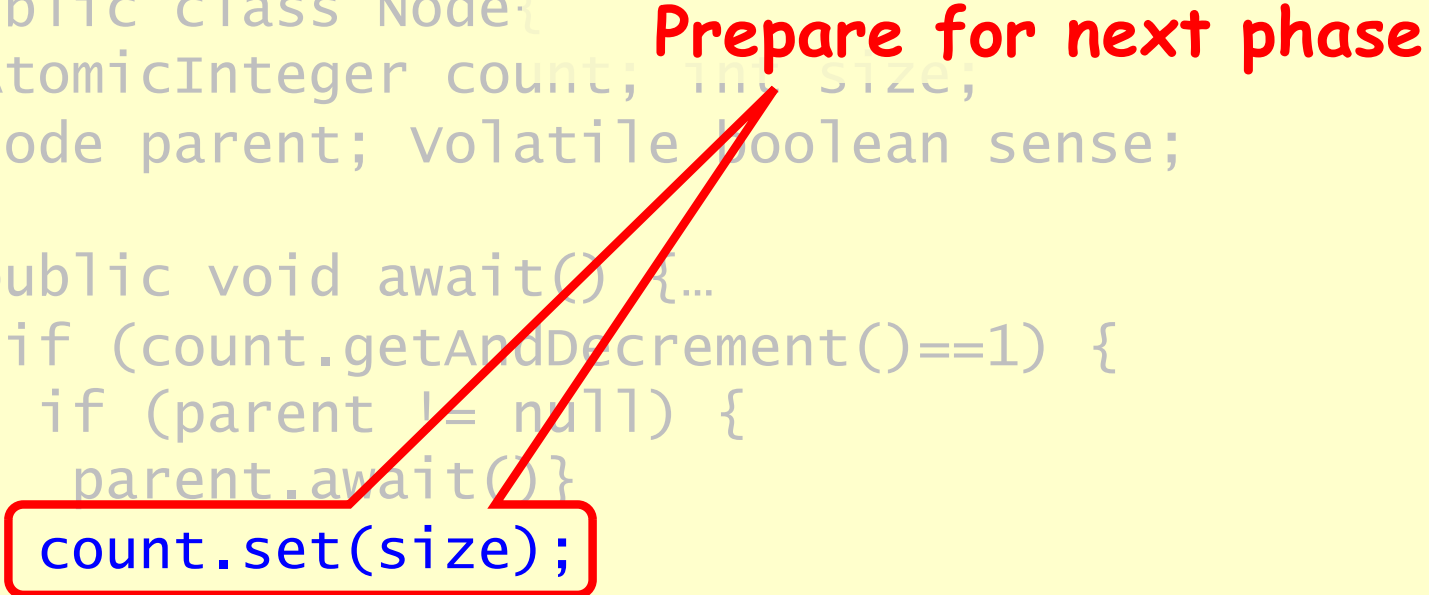


Combining Tree Barrier

```
public class Node{
    AtomicInteger count; int size;
    Node parent; volatile boolean sense;

    public void await() {...
        if (count.getAndDecrement()==1) {
            if (parent != null) {
                parent.await();
                count.set(size);
                sense = mySense
            } else {
                while (sense != mySense) {}
            }
        }
    }
}
```

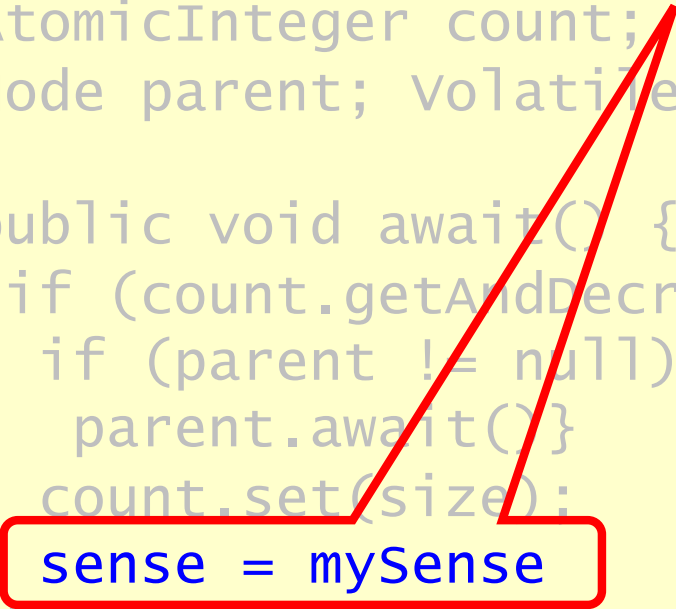
Prepare for next phase



Combining Tree Barrier

```
public class Node{ Notify others at this node
    AtomicInteger count; int size;
    Node parent; volatile boolean sense;

    public void await() {...
        if (count.getAndDecrement()==1) {
            if (parent != null) {
                parent.await();
                count.set(size);
                sense = mySense
            } else {
                while (sense != mySense) {}
            }
        }
    }
}
```



Combining Tree Barrier

```
public class Node{
    AtomicInteger count; int size;
    Node parent; volatile boolean sense;

    public void await() {...
        if (count.getAndDecrement()==1) {
            if (parent != null) {
                parent.await();
                count.set(size);
                sense = mySense;
            } else {
                while (sense != mySense) {}
            }
        }
    }
}
```

I'm not last, so wait for notification

A red line starts from the 'while (sense != mySense) {}' line, goes up and to the right, then turns left to point at the text 'I'm not last, so wait for notification'. Another red line goes down and to the left from the same point, then turns left to point at the 'while' loop.

Combining Tree Barrier

- No sequential bottleneck
 - ❑ Parallel getAndDecrement() calls
- Low memory contention
 - ❑ Same reason
- Cache behavior
 - ❑ Local spinning on bus-based architecture
 - ❑ Not so good for NUMA